

CS 426: Computer Security

Fall 2025

Homework 3: Advanced Buffer Overflow Attacks

Due Sept. 18th, 2025 @ 11:59 PM

1 Overview

The goal of this assignment is to help you gain hands-on experience on buffer overflow attacks. Before you start, you might want to go over Aleph One's seminal article *Smashing the Stack for Fun and Profit* available at this link https://inst.eecs.berkeley.edu/~cs161/fa08/papers/stack_smashing.pdf [One96]. You can also review the example we saw in class (or watch the recording of) on Sept. 4th, Sept. 9th and Sept. 11th.

2 Environment Setup

Use the provided Kali Linux VM for this assignment. Buffer overflow exploitation is sometimes delicate, and their success depends on specific details of the target system, so you will be using the same Kali Linux VM used in your previous homework to develop and test your attacks. This will also promote a consistent experience for everyone in the class and facilitate the grading process. We have slightly tweaked the configuration to disable some of the security features that are commonly used in the wild but would complicate our buffer overflow attacks. We will use this precise configuration to grade your submissions, so do not use your own VM instead. You must download the targets into a folder owned by the VM; running targets in a shared folder might cause incorrect behavior. All materials for this question are available via Brightspace.

If you have not installed the VM, the image (`cs426-kali.ova`) file can be downloaded from this link: <https://app.box.com/s/x8hcgawqvdy9qipwca54sw7gxmsp3tgi>.¹ See the Homework 1 document for the installation details.

You must not attack anyone else's system without authorization! This assignment asks you to develop attacks and test them in a virtual machine you control. Attempting the same kinds of attacks on others' systems without authorization is prohibited by law and university policies and may result in *fines, expulsion, and jail time*. Per the course ethics policy, you are required to respect the privacy and property rights of others at all times, *or else you will fail the course*.

¹This file is quite large (2 GB), so we recommend that you do not download it with your mobile data.

3 Assignment Setup

Files for the assignment are provided in Brightspace. Run the VM, download the file `hw3.tar.gz`, and place it under `/home/cs426/`. Then, open the terminal and navigate to `/home/cs426/` to unpack the package and compile the vulnerables. For these steps, run the commands below from your terminal.

1. Inside Kali VM, download the `hw3.tar.gz` package and unpack it.

```
cd /home/cs426/
wget -O hw3.tar.gz \
    https://app.box.com/shared/static/bpmjru026fyw01aeqbdiipzi9rw5zyyud.gz
tar zxvf hw3.tar.gz
cd hw3
```

2. Compile the randomized target binaries. First, set a cookie using your Purdue Career username (e.g., `./setcookie ling102`). Then, generate the randomized target source codes based on the cookie, and compile them into binaries:

```
./setcookie <username>          # replace <username> with your Purdue username
make generate
echo "cs426" | sudo -S make
```

Now you can start attacking the target binaries! Here are some requirements and hints that will help you in this assignment:

- **No specialized attack tools!** You may not use special-purpose tools meant for testing the security or exploiting vulnerabilities. You must complete the problem using only general purpose tools, such as `gdb`.
- **GDB:** You will find the GDB debugger extremely useful for finding buffer overflow exploits. Make use of your prior experience from Homework 1, and do not be afraid to experiment! This quick reference may also be useful: <https://users.ece.utexas.edu/~adnan/gdb-refcard.pdf>.
- **x86 Assembly:** There are many good references for x86 assembly language but note that this assignment targets the 32-bit x86 ISA. The stack is organized differently in x86 and x86_64. If you are reading any online documentation, ensure that it is based on the x86 architecture, e.g., <https://cs.dartmouth.edu/~sergey/cs258/tiny-guide-to-x86-assembly.pdf>.

3.1 Miscellaneous

You may find it helpful to write an additional program or script that will assist in figuring out the correct input to exploit any of the vulnerabilities. Note that this can be separate from the script/program you will write to exploit the vulnerabilities. If you are unsure about what is acceptable for this assignment, please do not hesitate to contact us. Be sure to turn in the code for any additional program or script that you write.

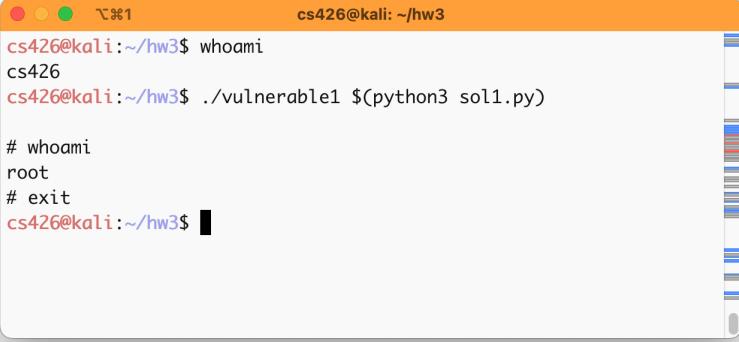
4 Targets

This assignment requires you to attack three vulnerable C programs by identifying and exploiting their vulnerabilities. You will do this by writing three attack scripts, one for each target program. These scripts will provide malicious inputs or command-line arguments to exploit the vulnerabilities. You must use the default Python3 version provided in the Kali VM. Do not change the Python3 version or modify the script links.

IMPORTANT: Before attacks, make sure you have generated the `cookie` with your Purdue username, and compiled your target binaries with your `cookie` as in Section 3. Otherwise, we cannot grade your work. You can verify your cookie by uploading it to Gradescope under Homework 3 (Code).

4.1 vulnerable1: Overwriting the return address indirectly (10 points)

Similar to Tasks 3 and 4 in Homework 2, your goal is to gain root privilege as a normal user by exploiting buffer overflow vulnerabilities in the target programs. These target programs are owned by the `root` user and have the `suid` bit set. Under normal execution, these programs read in a cmdline argument and user input, perform some operations, and exit. Your task is to provide a cmdline argument or user input that causes the program to spawn a root shell, as in Figure 1.



The screenshot shows a terminal window titled 'Terminal' with the command line 'cs426@kali: ~/hw3\$'. The user runs 'whoami' and gets 'cs426'. Then they run './vulnerable1 \$(python3 sol1.py)'. The output shows the user has become root, with the command '# whoami' outputting 'root' and '# exit' being run. This indicates a successful exploit.

Figure 1: Successful run.

However, this time, the buffer overflow is restricted and cannot directly overwrite the return address. Create a Python3 program named `sol1.py` that prints a line to be used as the command-line argument to the target. Test your program with the command line:

```
./vulnerable1 $(python3 sol1.py)
```

On success, you will see a shell prompt (#), and running `whoami` will output "root" as in Figure 1.

Do not try to create a solution that depends on you manually setting environment variables. You cannot assume that the grader will run your solution with the same environment variables that you have set. Please make sure that your program exits gracefully when you close the spawned shell as well (it should not segfault).

HINTS: You should use the shellcode we have provided in `shellcode.py`². You can simply import it in your solution script (`from shellcode import shellcode`). You are not allowed to use any other (custom) shellcode in any of the exercises. Successfully placing this shellcode in memory and setting the instruction pointer to the beginning of the shellcode (e.g., by returning or jumping to it) will open a shell.

4.2 vulnerable2: Variable stack position (10 points)

Address Space Layout Randomization (ASLR) is a security defense that makes buffer overflow attacks more difficult. Instead of the stack being in a fixed location, ASLR randomly changes its starting address for each program execution.

For this attack, the stack's position is randomly offset by 0 to 255 bytes every time the program runs. Your task is to create an input that can successfully exploit the buffer overflow and **always** open a root shell, even with this randomization. Create a Python3 program named `sol2.py` that prints a line to be used as the command-line argument to the target. Test your program with the command line:

```
./vulnerable2 $(python3 sol2.py)
```

On success, you will see a shell prompt (#), and running `whoami` will output "root" as in Figure 1.

HINTS: For this assignment, your attack script must reliably overcome stack address randomization to successfully open a root shell every time it is executed. If your script fails even once, you will lose points.

4.3 vulnerable3: Bypassing DEP and Stack Canary (10 points)

Besides ASLR, many other defenses exist, including data execution prevention (DEP) and stack canaries. With DEP, data-holding memory regions are marked as non-executable, preventing execution of shell code from the stack. With stack canaries, random numbers are placed between local variables and the saved frame pointer (as well as the return address). If an attacker tries to overflow a local buffer and overwrite the frame pointer or return address, the program will detect the corruption of the stack canary and terminates execution.

For this attack, you need to bypass both defenses and gain root shell access through a buffer overflow attack. Normally, the program takes in the filename to operate on from the command line, and performs read/write operations based on instructions from users. You need to discover the vulnerability in this program, and carefully design a buffer overflow attack to bypass the two defenses. To accomplish the attack, create two text files:

- `sol3_input.txt`: This file contains the data that the program will read from and save to.
- `sol3_commands.txt`: This file contains all instructions for the program, telling it what operations to perform (e.g., read, write, save, quit).

²For reference, this shellcode is based on Aleph One's shellcode in [One96].

To test your attack manually, run the following command:

```
./vulnerable3 sol3_input.txt
```

and type each of your instructions by hand. To test your attack in an automatic way, run the following command:

```
cat sol3_commands.txt - | ./vulnerable3 sol3_input.txt
```

On success, you will enter a shell and be able to run shell commands, even though there's no shell prompt (#). Running `whoami` will output "root" as in Figure 1.

HINTS:

1. Due to some limitations, the Kali VM does not fully enforce DEP. However, you should pretend that it does. Therefore, do not attempt to inject shellcode and execute it as previous attacks, otherwise you will lose points.
2. To successfully gain root access with your exploit, you must first switch the system's default shell from `bash/dash` to `zsh`. This is because newer versions of `bash` and `dash` contain security features that prevent this type of exploit from working.

Run these commands to link `/bin/sh` to `zsh`:

```
sudo rm /bin/sh
sudo ln -s /bin/zsh /bin/sh
```

With the default shell `bash/dash`, you can only spawn a shell without root privileges, due to their advanced protections against privilege escalation attacks. When they detect they are being run in a `setuid` process (like your exploit), they immediately switch from the effective user ID (root) back to the real user ID (cs426), effectively preventing the privilege escalation. `zsh` does not have this specific protection, which is why the exploit works with it.

More details can be found on the man page for the `system` call (<https://linux.die.net/man/3/system>) and `bash` (<https://linux.die.net/man/1/bash>).

3. If you need ideas on how to bypass DEP and stack canary, you can review the examples we saw in class (or watch the recording of) on Tue Sept. 9 and Thu Sept. 11. Essentially, to bypass the stack canary defense, you must avoid corrupting the stack canaries. To bypass the DEP, you need to find a way to spawn a shell without executing shellcode on the stack. A common method is to call the `system()` function with `"/bin/sh"` as the argument.

After your attack, the program must exit gracefully without a segmentation fault. As mentioned in class, this can be achieved by providing a valid return address for the function you hijack.

Be careful when preparing the command string (e.g. `" /bin/sh"`) for the `system()` call. You can either find the string in the program's memory or inject it on your own. If you inject the string onto the stack, ensure it is placed above the hijacked function's stack frame to prevent it from being overwritten.

5 Autograding

We will grade your attack scripts with an autograder (not the one on Gradescope), under the same Kali VM with the same commands as in Section 4. Your attack scripts should **NOT** cause the target programs to print any extra characters or white space, or crash with a segmentation fault. The screenshot below shows the expected outputs for successful attacks.



The screenshot shows a terminal window titled "cs426@kali: ~/hw3". The user has run three commands:

- `./vulnerable1 $(python3 sol1.py)` which prints "# whoami" followed by "root" and "# exit".
- `./vulnerable2 $(python3 sol2.py)` which prints "# whoami" followed by "root" and "# exit".
- `cat sol3_commands.txt - | ./vulnerable3 sol3_input.txt` which prints a series of command-line options for reading and writing memory, followed by "Exiting application", "whoami", "root", and "exit".

6 Deliverable

6.1 What and How to Submit

You need to submit your attack scripts and a report on **Gradescope**, separately.

The PDF report should be submitted on Gradescope under *Homework 3: Advanced Buffer Overflow Attacks(Report)*. See the next subsection ([Section 6.2](#)) for what to include in the report. Please assign pages to each attack on Gradescope.

The attack scripts should be submitted on Gradescope under *Homework 3: Advanced Buffer Overflow Attacks(Code)*. You can use the `prepare_submission.sh` provided to install `zip` and compress all files needed for submission into `submission.zip` as below:

```
cd /home/hw3; chmod +x prepare_submission.sh; ./prepare_submission.sh
```

You will get a warning if there is a file missing. You can update the FILES list inside the script to include other files you want to submit.

We will perform only basic checks on Gradescope (e.g. file naming and cookie value check). After the deadline, we will test your attack scripts using another autograder as specified in [Section 5](#).

File	Description	Submission
<code>report.pdf</code>	Your write-up answering the questions in Section 6.2 .	<code>pdf</code> on Gradescope [Homework 3 (Report)]
<code>cookie</code>	Your cookie file used for personalizing the <code>vulnerable</code> s (Section 3)	
<code>sol1.py</code>	Your exploit file attacking <code>vulnerable1</code> (Section 4.1)	
<code>sol2.py</code>	Your exploit file attacking <code>vulnerable2</code> (Section 4.2)	
<code>sol3_input.txt</code>	Your text file attacking <code>vulnerable3</code> (Section 4.3)	
<code>sol3_commands.txt</code>	Your commands attacking <code>vulnerable3</code> (Section 4.3)	
<i>other files</i>	Any additional files or programs created to support your efforts (Section 3.1)	All remaining files should be submitted on Gradescope [Homework 3 (Code)]

Table 1: List of required submission files

6.2 Homework Report Instruction

Your submission should include a report/write-up containing enough instructions and technical details to prove the **validity** and **reproducibility** of your work and demonstrate your understanding of each problem and your solution. More specifically, your document should contain detailed answers to the following questions for all `vulnerable`'s:

- (a) What does the program do? Describe the behavior of the program.
- (b) Why is this program vulnerable (living up to its name)? Identify the vulnerabilities this program has and explain why/how exactly they are vulnerabilities.

- (c) What could happen to the machine running this program? Describe the implication and possible outcome of the vulnerabilities identified above.
- (d) How can you exploit the identified vulnerabilities? Describe your strategy for attacking this program. Figures like the stacks we drew during lectures are welcomed.
- (e) What was your code? Attach your code.³ Screenshots are also acceptable. Make sure to keep your code as concise as possible (e.g., remove all unnecessary lines), but leave a sufficient amount of comments so it shows that your code matches the strategy you illustrated in the previous part. If your code is incomplete, still leave what you have gotten so far and describe the direction you were heading and what your plan was to reach the goal.
- (f) What proves that your attack was successful? Attach the screenshot(s) of the result of your attack showing that your attack indeed worked and achieved our goal. If your attack was not successful, still leave a screenshot of the result of your code and describe your debugging experience: What was the failure symptom? Why do you think your code was failing and why do you think it is a reasonable hypothesis? How do you think you could (fix that bug and) make it work?
- (g) If you were the author of the program (`vulnerable n`), you would want to ‘update’ the program so your users do not suffer the same vulnerabilities. How would you remediate/eliminate these vulnerabilities? That is, which line(s) would you replace (and with what)? If this cannot be fixed simply by replacing some lines, what other defense techniques would you use? Please explain with details.

Please be aware: If your report lacks details on many fronts, you may receive a grade close to the grade item for **no attempt or equivalent** on the rubric, EVEN IF your attack is correct. It is also your responsibility to make sure your responses are clean and organized enough to be readable. Similar to Homework 1, meaningless or frivolous responses such as “*because it is correct*”, “*trivial*”, and “*left as an exercise to the graders*” will be considered completely invalid. In particular, for Part (g), please suggest a fix instead of leaving an unproductive solution such as “*just do not use this program in the first place*” or a copy-paste (or equivalent) of lecture slides.

On the last page of your report, include a **collaboration statement**. It can simply be a list of people you communicated with and the online resources you have used for this homework. If you have finished this assignment entirely on your own AND did not consult any external resources, you MUST indicate that on your collaboration statement. This is a part of your assignment. **Failure to provide one could result in a deduction of points.**

³If you are using L^AT_EX, you may find `lstlisting` environment from the `listings` package useful.

7 Grading and Policy

7.1 Grade Breakdown

Each attack is worth 10 points, with 5 points for the code and 5 points for the corresponding write-up in your report. In total, this homework is worth 4% of your final grade.

As mentioned in Section 5, the code will be graded with an autograder. The write-up will be graded based on the rubrics in Section 6.2, with a focus on your understanding of the vulnerabilities, your attack strategies, and your proposed mitigations.

Even if your attacks failed, you can still earn half of the points through your write-up, so don't give up! If your exploit codes keep failing and you are unable to fix them, submit your best attempt along with a detailed explanation in your report. Clearly describe your approach, what you were trying to accomplish, and where you believe your code is failing. Depending on your reasoning and how close you are to the correct solution, you may still receive partial credit.

7.2 Important Rules and Reminders

This assignment MUST be completed using only `gdb` and `python3`, along with necessary general-purpose tools such as `gcc` and `make`. DO NOT use other special-purpose exploit tools such as `ghidra`. No points will be given to a solution that utilizes such tools. For all attacks, DO NOT try to create a solution that depends on you manually setting environment variables.

You are more than welcome to use any external resources including the ones available online besides the ones provided to you already (e.g., [One96]). However, the academic integrity policy still holds. **Do not ask someone else to do your work for you** or copy and paste the questions on websites to find answers. **Use your resources reasonably, and do not cross the line**. See the **Academic Integrity** section of the syllabus (course webpage) for more details.

References

- [One96] Aleph One, *Smashing the stack for fun and profit*, 1996, Available Online: https://inst.eecs.berkeley.edu/~cs161/fa08/papers/stack_smashing.pdf.